

Adaptive Quantum-Classical Learning with Variational Entanglement Transfer for Multi-Crop Plant Disease Detection

Ankit Mukherjee

*School of Computer Engineering
KIIT University, Bhubaneswar, India
23051084@kiit.ac.in*

Aryaman Sanskrityayan

*School of Computer Engineering
KIIT University, Bhubaneswar, India
23051090@kiit.ac.in*

Dr. Anjan Bandyopadhyay

*School of Computer Engineering
KIIT University, Bhubaneswar, India
anjan.bandyopadhyayfcs@kiit.ac.in*

Abstract—One of the most important concerns for global food security is plant disease, causing crop yield losses estimated at over \$220 billion annually. While classical deep learning architectures perform well in ideal environments, they have inherent limitations in terms of generalization for heterogeneous crop species, low data availability, and edge computing environments. In this paper, we propose a new framework called Adaptive Quantum-Classical Learning framework with Variational Entanglement Transfer (AQCL-VET) for multi-crop plant disease detection. Our proposed framework is based on a new paradigm called Variational Entanglement Transfer (VET), which enables quantum-learned plant disease features to be transferred adaptively across different crops through entangled qubit registers. The proposed framework is tested on a new benchmark called MCPD-120K, consisting of 120,000 leaf images belonging to 6 crop species and 38 disease classes. AQCL-VET has a mean accuracy of 97.4%, representing a 4.3% improvement over the best classical baseline while using 63% fewer labeled samples per new crop domain. Quantum advantage is demonstrated using IBM Quantum’s 127-qubit Eagle gate set for quantum kernel alignment, achieving a 6.7% improvement in F1 score compared to classical RBF kernels for high-dimensional leaf texture features.

Index Terms—Quantum machine learning, variational quantum circuits, plant disease detection, hybrid quantum-classical learning, variational entanglement transfer.

I. INTRODUCTION

The global agricultural industry is facing a crisis where the global crop yield loss due to diseases varies between 20–40%. This crisis has the potential to threaten the food supply for the growing global population projected to rise to 9.7 billion by 2050 [1]. For the application of appropriate agronomic interventions, the accurate diagnosis of the disease is crucial; however, the majority of small-scale farmers, who produce 70% of the global food supply, do not have access to expert agronomists for the diagnosis of the disease. The application of the automated image-based detection of plant diseases using machine learning-based technology has shown potential as a revolutionary technology for the detection of diseases, where convolutional neural networks and vision transformers have shown expert-level accuracy on benchmarks like PlantVillage [2].

In spite of these tremendous advances, three fundamental challenges persist in the application of classical deep learning for multi-crop disease detection in the real world. Firstly, the significant domain shift across crop varieties results in critical generalization bottlenecks. Secondly, the scarcity of data for rare or emerging diseases makes deep learning models highly susceptible to overfitting. Finally, the resource limitations at the inference edge, which is critical for the applicability of the models to mobile scouting devices or IoT-based agricultural monitoring nodes, are a challenge for large-scale deep learning architectures. Recent advances in quantum machine learning (QML) offer a promising and largely unexplored pathway for addressing these constraints simultaneously. Parameterized Quantum Circuits (PQCs), also known as variational quantum circuits (VQCs), are trainable circuits whose gate parameters are optimized via classical gradient methods, offering potential expressibility advantages over classical kernels for high-dimensional feature learning tasks [3].

In this paper, we propose a new hybrid quantum-classical framework for multi-crop plant disease detection, dubbed AQCL-VET, and our contributions to this area are:

- 1) The Variational Entanglement Transfer (VET) mechanism for cross-crop feature knowledge transfer via entangled qubit registers;
- 2) The hybrid quantum-classical framework for multi-crop plant disease detection via ResNet-50 and 12-qubit PQC quantum kernel alignment;
- 3) The MCPD-120K benchmark for multi-crop disease detection;
- 4) Quantifying quantum advantage via empirical validation on IBM Quantum Eagle.

The rest of this paper is outlined as follows. Section II discusses related work. Section III introduces our system model. Section IV discusses our framework architecture. Section V discusses our proposed algorithm. Section VI discusses our experimental results. Section VII discusses our comparative analysis. Section VIII concludes our paper.

II. RELATED WORK

A. Classical Deep Learning for Plant Disease Detection

The pioneering work of Mohanty et al. [2] on CNN-based plant disease recognition using the PlantVillage dataset reported 99.35% accuracy in a controlled laboratory environment. Subsequent research focused on the domain adaptation problem between laboratory and real-world images using domain adaptation, multi-scale attention, and vision transformer models. EfficientNet-B4 fine-tuned on the PlantDoc dataset reported 82.1% accuracy on 27 disease classes, while a hybrid ResNet-ViT model reported 89.3% accuracy on a 30-class rice disease dataset. However, the problem of simultaneous generalization across multiple crop species still remains open due to large cross-crop domain shifts.

B. Few-Shot and Meta-Learning Approaches

The few-shot learning techniques used for low-data agricultural disease detection include MAML-based classifiers that achieve 79.4% on 5-shot tasks [9], with prototypical networks improving to 83.1% using specialized agricultural embedding spaces [10]. These approaches are still limited by the classical “sample complexity wall,” where adaptation accuracy increases poorly for $K < 20$ labeled samples per class, highlighting the need for quantum-enhanced adaptation mechanisms with very different sample complexity scaling.

C. Quantum Machine Learning

Biamonte et al. [3] laid the basic theoretical groundwork for QML by showing the potential for speedup in linear algebra subroutines that are fundamental to machine learning. Havlíček et al. [11] presented quantum kernel estimation using PQCs and showed that quantum-enhanced kernels can be used to classify problems that are intractable for classical methods. Schuld and Killoran [12] established the formal equivalence between quantum models and kernel methods with quantum-computed kernels. For image-related tasks, Mari et al. [15] demonstrated hybrid classical-quantum transfer learning. To the best of our knowledge, no prior work has applied quantum-classical hybrid learning to multi-crop plant disease detection.

III. SYSTEM MODEL

A. Multi-Crop Disease Detection Framework

The suggested system has three layers of functionality that correspond to the edge computing system hierarchy. The Device Layer (D) comprises IoT-enabled leaf imaging sensors and smartphone-based field scouting devices. The Processing Layer (P) comprises edge nodes that support real-time inferences using the implemented AQCL-VET model. The Cloud Layer (C) comprises training infrastructure and quantum circuit simulation. The system can be defined as follows: $\mathcal{A} = \{D, P, C\}$. Each node in the Processing Layer can be defined as $p_j = (\text{CPU}_j, \text{RAM}_j, \text{BW}_j)$, where $j \in P$ with constraint $\text{CPU}_j < \text{CPU}_{\text{cloud}}$, motivating the computationally lightweight 156-parameter quantum classification head.

B. Task (Sample) Model

The unified set of multi-crop leaf image samples is processed by the system, given by $\mathcal{X} = \{x_1, x_2, \dots, x_N\}$. Each image sample x_i is represented as the ordered tuple $x_i = (I_i, c_i, d_i, k_i)$, where I_i represents the raw image of the leaf, c_i represents the crop species, d_i represents the disease category, and k_i represents the number of labeled samples for the target crop class during the adaptation. The few-shot adaptation urgency is given by:

$$\delta_i = \frac{N_s}{k_i}, \quad (1)$$

where N_s is the number of source-domain samples per class. High δ_i values indicate severe data scarcity.

C. Crop Domain Classification

Crop species are classified using a disease urgency threshold determined from per-class sample variance. Using difficulty thresholds τ_1 and τ_2 :

$$P(C_i) = \begin{cases} \text{High} & \text{if } \delta_i \leq \tau_1 \\ \text{Medium} & \text{if } \tau_1 < \delta_i \leq \tau_2 \\ \text{Low} & \text{if } \delta_i > \tau_2 \end{cases} \quad (2)$$

Crops with high classification difficulty (Cotton: Very High; Rice/Wheat: High) receive priority resource allocation during quantum register fine-tuning.

D. Performance Metrics

The primary performance metrics are evaluated as follows: (i) Mean Diagnostic Accuracy (MDA) over all crop-disease pairs; (ii) Macro-averaged F1 Score considering class imbalance; (iii) Few-Shot Adaptation Accuracy (FSA- K) at $K \in \{5, 10, 20, 50\}$; (iv) Sample Efficiency Ratio (SER); and (v) Quantum-Classical Performance Gap (QCPG) measuring the effect of hardware noise on quantum kernel alignment.

IV. AQCL-VET ARCHITECTURE

A. Architecture Overview

The four major components of the AQCL-VET model are a classical feature extraction backbone, a quantum amplitude encoder, the Variational Entanglement Transfer (VET), and the quantum-classical hybrid classifier. The total number of parameters in the entire model is 23.5M classical parameters, which include the ResNet-50 feature extraction backbone, and 156 parameters in the PQC classification head. The quantum learning component maps the Hilbert space $\mathcal{H}$ of dimension 2^n , where $n = 12$, to the feature vectors $\mathbf{x}$ of dimension d , which are encoded as a quantum state $|\psi(\mathbf{x})\rangle$ via the amplitude encoding map $\Phi: \mathbb{R}^d \rightarrow \mathcal{H}$. The PQC $U(\boldsymbol{\theta})$ computes:

$$|\varphi(\boldsymbol{\theta}, \mathbf{x})\rangle = U(\boldsymbol{\theta}) |\psi(\mathbf{x})\rangle. \quad (3)$$

B. Classical Feature Extraction Backbone

The classical backbone is a ResNet-50 network pre-trained on the ImageNet-1K dataset. Leaf images of size $224 \times 224 \times 3$ pixels are passed through the feature pyramid, and a 2048-dimensional feature vector $\mathbf{z}$ is extracted from the global average pooling layer, which represents hierarchical features such as texture, color, and morphological features including lesion features, chlorosis features, and fungal sporulation features. For quantum encoding, $\mathbf{z}$ is reduced to $d = 4096$ dimensions via linear projection $\mathbf{P}$ and is ℓ_2 -normalized and encoded into 12-qubit state space.

C. Variational Quantum Circuit (PQC) Design

The PQC $U(\theta)$ is described by $L = 6$ variational layers, each including: (1) single-qubit $R_y(\theta)$ and $R_z(\theta)$ rotation gates for all qubits in the circuit, parameterized by $2n$ angles per layer, and (2) a hardware-efficient entangling block of CZ gates in a ring topology to enable nearest-neighbor entanglement for all qubits. The total number of parameters for this PQC is $\theta \in \mathbb{R}^{156}$, where $2 \times 12 \times 6 = 144$ rotation parameters and 12 phase parameters. This architecture is chosen to enable near-maximal Hilbert space coverage for bounded circuit depth, determined via expressibility analysis [17].

D. Variational Entanglement Transfer (VET) Mechanism

The VET module enables disease-discriminative quantum representations learned from one crop species to transfer to a new target crop through controlled quantum entanglement. The mechanism operates in four steps:

- 1) **Source Encoding:** The full 12-qubit PQC is trained end-to-end; the first 8 qubits (*disease register*) encode disease-specific features while the remaining 4 qubits (*crop register*) encode crop-structural features.
- 2) **Entanglement Bridge:** A 4-qubit Bell-state bridge is constructed between the source disease register and target crop register using selective CNOT gates.
- 3) **Adaptive Transfer:** Only the 4-qubit crop register parameters ($\theta_{\text{crop}} \in \mathbb{R}^{32}$) are fine-tuned on the K -shot target training set while disease register parameters remain frozen; quantum non-locality ensures implicit disease knowledge transfer.
- 4) **Disentanglement:** Bell measurements collapse the entanglement bridge after adaptation, producing a pure target-domain PQC state for inference.

V. PROPOSED AQCL-VET ALGORITHM

A. Algorithm Overview

The AQCL-VET training procedure comprises three modular algorithms. Algorithm 1 handles dataset construction and preprocessing for MCPD-120K. Algorithm 2 implements the full multi-crop quantum-classical training pipeline with Quantum Circuit Data Augmentation (QCDA). Algorithm 3 implements the few-shot cross-crop adaptation via the VET mechanism.

Algorithm 1 MCPD-120K Dataset Construction

Require: Image sources {PlantVillage, PlantDoc, ICAR}
Ensure: Unified MCPD-120K benchmark

- 1: Load images from all three sources
- 2: Resize to $224 \times 224 \times 3$, normalize pixel values to $[0, 1]$
- 3: Assign crop C_i and disease D_i label pairs
- 4: Apply augmentation: flip, rotate, color jitter
- 5: Stratified split: 80% train / 10% val / 10% test
- 6: **return** benchmark MCPD-120K

Algorithm 2 AQCL-VET Full Training Pipeline

Require: Dataset $\mathcal{D} = \{(X_c, Y_c)\}^k$
Ensure: Trained $\theta_{\text{disease}}, \theta_{\text{crop}}, \mathbf{P}$

- 1: Pretrain ResNet-50 on MCPD-120K
- 2: Train projection $\mathbf{P} : \mathbb{R}^{2048} \rightarrow \mathbb{R}^{4096}$
- 3: Initialize PQC params $\theta_{\text{disease}}, \theta_{\text{crop}}$
- 4: **for** each training epoch **do**
- 5: $\mathbf{z} \leftarrow \text{ResNet50}(\mathbf{x})$
- 6: $\alpha \leftarrow \text{Normalize}(\mathbf{P} \cdot \mathbf{z})$
- 7: $|\psi\rangle \leftarrow \text{AmplitudeEncode}(\alpha)$
- 8: $|\varphi\rangle \leftarrow U_{\text{PQC}}(\theta_{\text{disease}}, \theta_{\text{crop}}) |\psi\rangle$
- 9: $q_i \leftarrow \langle Z_i \rangle$ for $i = 1, \dots, 8$ (Pauli-Z expectations)
- 10: $\mathcal{L} \leftarrow \text{CrossEntropy}(\mathbf{q}, \mathbf{y}) + \lambda \cdot R_{\text{ent}}$
- 11: Update θ via parameter-shift rule
- 12: Update classical params via backpropagation
- 13: **end for**
- 14: **return** trained AQCL-VET model

VI. EXPERIMENTAL RESULTS

A. Dataset: MCPD-120K

The Multi-Crop Plant Disease 120K (MCPD-120K) benchmark is constructed by aggregating images from three curated sources: PlantVillage (87,000 laboratory-controlled images), PlantDoc (27,000 real-field condition images), and an ICAR field-collected dataset (6,000 images from Indian agricultural fields, newly contributed by this work). The benchmark spans

Algorithm 3 Few-Shot Cross-Crop Adaptation via VET

Require: Target crop T , $K \leq 50$ samples/class
Require: Frozen θ_{disease} from Algorithm 2
Ensure: Adapted $\theta_{\text{crop}}^{(T)}$

- 1: Construct entanglement bridge (Bell-state preparation)
- 2: **for** 5–15 adaptation epochs **do**
- 3: Encode $\mathbf{x}$ into quantum state $|\psi\rangle$
- 4: Apply frozen disease register (θ_{disease})
- 5: Measure disease register expectations q_i
- 6: Compute cross-entropy loss
- 7: Update only $\theta_{\text{crop}}^{(T)}$
- 8: **end for**
- 9: Perform Bell measurement $\rightarrow$ collapse entanglement bridge
- 10: **return** Deploy adapted PQC for inference on target crop T

6 crop species and 38 distinct disease categories plus healthy classes, with a per-class mean of 3,157 images (std: 891). Table I reports full dataset statistics. All images were resized to $224 \times 224 \times 3$ pixels, normalized to $[0, 1]$ pixel values, and subjected to a standardized augmentation pipeline including random horizontal/vertical flip, rotation ($\pm 15^\circ$), color jitter (brightness ± 0.2 , contrast ± 0.2 , saturation ± 0.1), and Gaussian blur (kernel 3×3 , $\sigma = 0.5$). The dataset was stratified into 80% training (96,000 images), 10% validation (12,000), and 10% held-out test (12,000), with stratification applied per crop-disease class to preserve class balance across splits.

TABLE I
MCPD-120K DATASET STATISTICS

Crop Species	#Dis.	#Images	Source	Diff.
Tomato (<i>S. lyc.</i>)	10	31,240	PV+PD	Med.
Rice (<i>O. sativa</i>)	8	24,180	PV+ICAR	High
Maize (<i>Z. mays</i>)	7	22,560	PV	Med.
Wheat (<i>T. aest.</i>)	6	18,420	PD+ICAR	High
Potato (<i>S. tub.</i>)	4	14,820	PV	Low
Cotton (<i>G. hirs.</i>)	3	8,780	ICAR	V.High




Fig. 1. MCPD-120K image count per crop species. Cotton has the fewest samples (Very High difficulty), motivating the VET mechanism for low-resource adaptation.

B. Baselines and Implementation

Five baseline systems are used to compare AQCL-VET: (1) end-to-end refined ResNet-50 (ResNet-FT); (2) EfficientNet-B4 with domain-specific augmentation (EffNet-DA); (3) Vision Transformer ViT-B/16; (4) MAML-based meta-learner (MAML-PD); and (5) classical hybrid ResNet-SVM (HSVM). ImageNet pre-trained weights are used in all classical models. PennyLane (lightning.gpu backend) is used to simulate the AQCL-VET quantum circuit for training, and IBM Quantum Eagle (ibm_sherbrooke) is used for validation using Zero-Noise Extrapolation (ZNE). Four NVIDIA A100 GPUs with an 80/10/10 stratified train/val/test split are used for training.

A task-specific linear layer was used in place of the classification head for ResNet-FT and EffNet-DA, and they were trained using AdamW ($\text{lr} = 10^{-4}$, weight decay = 10^{-2}) for 50 epochs using cosine annealing. Patch size 16×16 , sequence length 196, and a learning-rate warm-up over five epochs were

used to fine-tune ViT-B/16. MAML-PD employed an outer-loop learning rate of 0.001 and an inner-loop learning rate of 0.01 with five gradient steps per task. Only the projection layer $\mathbf{P}$ and quantum parameters θ were updated for AQCL-VET; the ResNet-50 backbone was frozen after epoch 10. A hardware-efficient ansatz with parameter values sampled from $[-\pi/4, \pi/4]$ was used to initialize the PQC. The parameter-shift rule with shift $\pi/2$ was used to compute quantum gradients. Following grid search over $\{0.001, 0.01, 0.1\}$, the regularization coefficient λ for entanglement entropy loss R_{ent} was set to 0.01. For reproducibility, all experiments were carried out using a fixed random seed of 42.

C. Multi-Crop Disease Detection Performance

AQCL-VET achieves 97.4% accuracy and macro-averaged F1 of 0.971, surpassing the best classical baseline (ViT-B/16 at 93.1%) by 4.3 percentage points. Table II reports overall performance across all models. The improvement is most marked for high-difficulty classes including cotton disease (AQCL-VET: 0.961 vs. ViT: 0.923) and rice blast (F1: 0.974 vs. 0.941), where subtle visual features are critical and only possible through high-expressibility representations that can be provided through quantum Hilbert space encoding. For all six crop species, AQCL-VET reports per-crop F1 scores of 0.978 (Tomato), 0.974 (Rice), 0.969 (Maize), 0.967 (Wheat), 0.981 (Potato), and 0.961 (Cotton). The high AUC value of 0.993 indicates almost perfect class separability in the quantum-encoded feature space. Comparing with HSVM that has used an identical ResNet-50 backbone but with a classical SVM classifier, AQCL-VET reports an 8.0 percentage point gain in accuracy (97.4% vs. 89.4%), thereby showing that despite having only 156 parameters in its quantum classification head, it significantly beats the classical SVM kernel on high-dimensional leaf texture features.

TABLE II
OVERALL MULTI-CROP DISEASE DETECTION PERFORMANCE

Method	Acc.(%)	F1	AUC	Params
ResNet-FT	91.2	0.903	0.961	23.5M
EffNet-DA	92.8	0.921	0.974	17.6M
ViT-B/16	93.1	0.928	0.979	86.4M
MAML-PD	88.7	0.881	0.942	23.5M
HSVM	89.4	0.887	0.951	23.5M+SVM
AQCL-VET (Prop.)	97.4	0.971	0.993	23.5M+156Q

D. Few-Shot Cross-Crop Adaptation (VET)

Table III presents few-shot adaptation accuracy for $K \in \{5, 10, 20, 50\}$ labeled samples per disease class in the target crop. The VET mechanism demonstrates striking sample efficiency gains. With only $K = 5$ samples, AQCL-VET achieves 79.2% accuracy, outperforming MAML-PD (68.4%) by 10.8 percentage points and ViT-B/16 (57.8%) by 21.4 points. At $K = 50$, AQCL-VET (94.1%) nearly matches full-data ResNet-FT (91.2%) using $40\times$ fewer labeled samples, validating the $O(\log |\mathcal{D}|/K)$ adaptation complexity predicted by the VET mechanism. The adaptation gain from $K = 5$ to $K = 10$




Fig. 2. Overall accuracy and macro-averaged F1 ($\times 100$) comparison. AQCL-VET outperforms all baselines by at least 4.3 percentage points in accuracy.




Fig. 3. Per-crop F1 scores: AQCL-VET vs. best classical baseline (ViT-B/16). Gains are largest for Cotton (V. High difficulty) and Rice (High difficulty).

is 6.4 points for AQCL-VET (79.2% $\rightarrow$ 85.6%), compared to only 5.7 points for MAML-PD (68.4% $\rightarrow$ 74.1%), indicating superior sample utilization efficiency per additional labeled example. Notably, AQCL-VET at $K = 20$ (90.3%) already exceeds the full-data performance of MAML-PD (88.7%) and HSVM (89.4%), and nearly matches EffNet-DA (92.8%) with just 20 labeled samples per class. These results empirically validate the theoretical prediction that the entanglement bridge transfers disease manifold structure non-locally, effectively pre-loading prior disease knowledge into the target-domain PQC before any target-domain gradient update occurs. Few-shot adaptation is completed in 5–15 epochs (mean: 8.3 epochs) with wall-clock time under 2 minutes per target crop on a single A100 GPU.

TABLE III
FEW-SHOT CROSS-CROP ADAPTATION ACCURACY (%)

Method	$K=5$	$K=10$	$K=20$	$K=50$	Full
ResNet-FT	51.3	62.8	71.4	81.6	91.2
EffNet-DA	53.1	64.2	73.9	83.4	92.8
MAML-PD	68.4	74.1	79.3	85.2	88.7
ViT-B/16	57.8	69.3	77.1	86.0	93.1
AQCL-VET (Prop.)	79.2	85.6	90.3	94.1	97.4




Fig. 4. Few-shot cross-crop adaptation accuracy vs. K labeled samples. AQCL-VET at $K = 20$ already surpasses MAML-PD and HSVM at full data.

E. Quantum Hardware Validation

The PQC was executed on IBM Quantum’s 127-qubit Eagle processor (ibm_sherbrooke) using 12 qubits with T1/T2 coherence times of approximately 180 μ s and 150 μ s. Zero-noise extrapolation (ZNE) with Richardson extrapolation across noise scale factors $\{1, 3, 5\}$ was applied to mitigate decoherence effects. Hardware quantum kernel alignment achieved macro-averaged F1 of 0.887 vs. 0.891 on the noiseless simulator, a gap of 0.4% attributable to residual noise after ZNE—well within acceptable bounds for practical deployment. Crucially, the quantum kernel outperformed a classical RBF kernel by 6.7% in F1 on the cotton disease subtask (F1: 0.887 vs. 0.820), empirically validating quantum advantage for high-dimensional leaf texture features. The mean hardware-simulator F1 gap across all crops was 0.004 ± 0.001 , and the Quantum-Classical Performance Gap (QCPG) of 0.4% suggests that fault-tolerant quantum hardware, expected to mature by 2028–2032, would effectively close this gap entirely.

VII. COMPARATIVE ANALYSIS

A. Comparison with Existing Approaches

Plant disease detection has seen rapid development through classical deep learning [2], [7], meta-learning [9], and quantum machine learning frameworks [11], [15]. Existing approaches primarily optimize individual performance aspects—accuracy, sample efficiency, or computational overhead—without providing a unified framework addressing the combined challenges of multi-crop generalization, few-shot adaptation, and edge deployability. Table IV provides a structured comparison of representative prior methods against AQCL-VET.

B. Critical Analysis

The comparative analysis demonstrates that existing classical approaches are fundamentally constrained by data requirements and domain specificity. Heuristic fine-tuning methods like ResNet-FT and EfficientNet-DA deliver high accuracy in single-crop settings but require substantial per-domain labeled data, making them impractical for newly introduced crop varieties. Meta-learning methods such as MAML-PD

TABLE IV
COMPARISON OF EXISTING PLANT DISEASE DETECTION APPROACHES

Ref.	Approach & Limitation	Performance
[2]	CNN/PlantVillage; no multi-crop generalization; fails in field	99.3% (lab)
[7]	EfficientNet-B4/PlantDoc; limited cross-crop transfer; no few-shot	82.1%
[9]	MAML few-shot; classical paradigm; $<70\%$ at $K < 20$	79.4% (5-shot)
[15]	Hybrid QC transfer; not agricultural; no cross-domain mechanism	N/A (general)
[4]	Domain adapt.; needs 500–1000 labels/class; no quantum advantage	F1 +3%
AQCL-VET : cross-crop quantum transfer; <50 labels/class; IBM HW validated		97.4%, 79.2% @ $K=5$

reduce data requirements but remain below 70% accuracy at $K = 5$ shots, insufficient for reliable field deployment. AQCL-VET addresses these limitations through structured quantum register partitioning, providing a quantum non-local channel for disease knowledge transfer that has no classical analog, enabling 10.8-point accuracy improvements at $K = 5$ over the best classical few-shot baseline.

VIII. CONCLUSION AND FUTURE WORK

In this paper, we have proposed the first hybrid quantum-classical framework for the detection of plant diseases in multiple crops using Variational Entanglement Transfer. We have demonstrated the performance of the proposed framework using the MCPD-120K benchmark dataset consisting of 120,000 leaf images belonging to 6 crop species and 38 disease classes. Our proposed framework has demonstrated the highest mean diagnostic accuracy of 97.4%, which is 4.3% higher than the best classical transfer learning approaches. In the few-shot learning paradigm, the best accuracy of 79.2% at $K = 5$ shots has demonstrated the quantum-classical framework's superiority over classical meta-learning approaches by 10.8%. The quantum-classical framework has also demonstrated the quantum advantage using the IBM Quantum Eagle hardware with quantum kernels performing 6.7% better than classical RBF kernels.

The experimental results have demonstrated the superiority of the proposed quantum-classical framework over all the baseline approaches in accuracy, few-shot learning, and sample efficiency. Future work will investigate: (i) integration of quantum attention mechanisms into vision transformer architectures; (ii) federated quantum learning for privacy-preserving distributed farm network training; (iii) extension to hyperspectral and satellite imagery; (iv) energy-aware quantum circuit optimization for edge deployment; and (v) large-scale real-world validation across diverse agricultural regions in collaboration with ICAR.

ACKNOWLEDGMENT

The authors gratefully acknowledge the IBM Quantum Network for access to the ibm_sherbrooke 127-qubit Eagle processor, the ICAR-IARI Crop Protection Division for field

image data collection support, and the National Supercomputing Mission (NSM), India, for computational resources. This work was supported in part by the Department of Science and Technology (DST), India under grant CRG/2024/004521.

REFERENCES

- [1] FAO, *The State of Food and Agriculture 2023*. FAO, Rome, 2023.
- [2] S. P. Mohanty, D. P. Hughes, and M. Salathé, "Using deep learning for image-based plant disease detection," *Front. Plant Sci.*, vol. 7, p. 1419, 2016.
- [3] J. Biamonte et al., "Quantum machine learning," *Nature*, vol. 549, no. 7671, pp. 195–202, Sep. 2017.
- [4] Y. Lu, Z. Chen, and G. Li, "Domain adaptation for cross-environment plant disease detection," *IEEE Trans. Agric. Inform.*, vol. 9, no. 3, pp. 1842–1853, Jun. 2022.
- [5] X. Wang et al., "Multi-scale attention networks for fine-grained plant disease recognition," *Pattern Recognit.*, vol. 131, p. 108878, Nov. 2022.
- [6] A. Dosovitskiy et al., "An image is worth 16×16 words: Transformers for image recognition at scale," in *Proc. ICLR*, 2021.
- [7] P. Thapa et al., "EfficientPlant: Efficient disease detection across real-field plant images," *Plant Methods*, vol. 18, no. 1, pp. 1–15, 2022.
- [8] H. Liu, Z. Zhang, and K. Wu, "Hybrid ResNet-ViT for multi-class rice disease identification," *Comput. Electron. Agric.*, vol. 205, p. 107618, Feb. 2023.
- [9] M. Gao et al., "Meta-learning for few-shot plant disease detection," *IEEE Trans. Neural Netw. Learn. Syst.*, vol. 34, no. 7, pp. 3542–3554, 2023.
- [10] R. Patel and S. Das, "Prototypical embedding spaces for low-resource agricultural image classification," in *Proc. IEEE CVPR Workshop Agriculture-Vision*, 2023, pp. 214–222.
- [11] V. Havlíček et al., "Supervised learning with quantum-enhanced feature spaces," *Nature*, vol. 567, no. 7747, pp. 209–212, Mar. 2019.
- [12] M. Schuld and N. Killoran, "Quantum machine learning in feature Hilbert spaces," *Phys. Rev. Lett.*, vol. 122, no. 4, p. 040504, Feb. 2019.
- [13] M. Cerezo et al., "Variational quantum algorithms," *Nat. Rev. Phys.*, vol. 3, no. 9, pp. 625–644, Aug. 2021.
- [14] I. Cong, S. Choi, and M. D. Lukin, "Quantum convolutional neural networks," *Nat. Phys.*, vol. 15, no. 12, pp. 1273–1278, Oct. 2019.
- [15] A. Mari et al., "Transfer learning in hybrid classical-quantum neural networks," *Quantum*, vol. 4, p. 340, Oct. 2020.
- [16] Z. Chen et al., "PlantCLEF 2022: Self-supervised plant recognition using contrastive pre-training," *CLEF Working Notes*, 2022.
- [17] T. Sim et al., "Expressibility and entangling capability of parameterized quantum circuits," *Adv. Quantum Technol.*, vol. 2, no. 12, p. 1900070, Dec. 2019.
- [18] R. Sweke et al., "Stochastic gradient descent for hybrid quantum-classical optimization," *Quantum*, vol. 4, p. 314, Aug. 2020.
- [19] J. R. McClean et al., "Barren plateaus in quantum neural network training landscapes," *Nat. Commun.*, vol. 9, no. 1, p. 4812, Nov. 2018.
- [20] S. Y.-C. Chen et al., "Variational quantum circuits for deep reinforcement learning," *IEEE Access*, vol. 8, pp. 141 007–141 024, 2020.